

Machine Learning

SICCS-Oxford 2021

Dr Thomas Robinson, Durham University

June 2021

Hello!

Today's workshop:

- ▶ 1 hr 15 min lecture
- ▶ 15 minute break/Q&A
- ▶ 1.5 hour coding walkthrough on constructing neural networks in R

Caveat: three hours is not a lot of time!

- ▶ Introduce where I think ML is most useful in social sciences
- ▶ Equip you with some fundamental tools that can be applied across:
 - ▶ Contexts
 - ▶ Data sources
 - ▶ Algorithms

Lecture content

Goals are threefold:

1. Brief overview of machine learning

- ▶ What is ML?
- ▶ Prediction problems
- ▶ Bias-variance tradeoff

2. Building basic neural networks

- ▶ Highly flexible, "engineering-grade" ML method
- ▶ Now easily implementable in R

What is machine learning?

(Machine) learning and statistics

ML is a vague term:

*“Machine learning is a subfield of **computer science** that is concerned with building **algorithms** which, to be useful, rely on a collection of examples of some phenomenon. . . the process of solving a practical problem by 1) gathering a dataset, and 2) algorithmically building a statistical model based on that dataset.” – Burkov 2019*

To me, ML is defined by:

1. “Computationally-intensive” methods
2. Where researchers underspecify the relationship between variables
3. And allow the computer to search for (or learn) these relationships

Machine learning

Expectation: I need a \$1m super computer

Reality: It runs in minutes on a personal computer



Figure 1: Google: Tensor Processing Unit server rack

Why machine learning?

Machine learning can be:

- ▶ Powerful
- ▶ Flexible
- ▶ Reduce the burden on the researcher

It helps solve lots of **prediction problems** and can assist in **inference problems** too

But ML is not a panacea!

- ▶ ML cannot solve problems of poor research design
- ▶ And can introduce its own issues

Twitter apologises for 'racist' image-cropping algorithm

Users **highlight examples of feature automatically focusing on white faces over black ones**

Prediction and inference

Consider the following linear model:

$$\hat{y}_i = \hat{\beta}_0 + \hat{\beta}_1 x_{1i}$$

- ▶ Inference is concerned with estimating the size/direction of the relationship between variables ($\hat{\beta}$ problems)
- ▶ Prediction is concerned with estimating some outcome, using the relationships between variables (\hat{y} problems)

These two facets are clearly connected:

- ▶ If we know the size/direction of the relationships, we can predict the outcome
- ▶ But we rarely know (or even pretend to know) the true model
- ▶ Sometimes we can get good at \hat{y} problems without knowing $\hat{\beta}$

There are \hat{X} problems too

We can also think about where the prediction problem lies:

- ▶ \hat{y} problems are about the dependent variable
 - ▶ To predict an election winner...
 - ▶ ... or the probability of revolution...
 - ▶ ... or the weather tomorrow
 - ▶ These are not necessarily inferential problems
- ▶ \hat{X} problems are about independent variables
 - ▶ Dimensions of interest that may be important to our theory...
 - ▶ ... but which are not directly observable (i.e. latent)
 - ▶ We want to make predictions over \mathbf{X} so we can test an inferential theory about the relationship between X and y

Bias and variance


Bias is a feature of the estimator:

- ▶ $\text{Bias}_\beta = (\mathbb{E}[\hat{\beta}] - \beta)$
- ▶ With OLS under Gauss Markov assumptions, $(\mathbb{E}[\hat{\beta}] - \beta) = 0$

Variance occurs due to resampling from the population:

- ▶ Parameter estimates change (slightly) as we re-estimate the model with new data
- ▶ $\mathbb{V}_{\hat{\beta}} = \mathbb{E}[(\mathbb{E}[\hat{\beta}] - \hat{\beta})^2]$
- ▶ The average distance between a particular parameter estimate and the mean of parameter estimates over multiple samples

Visualising bias and variance



Bias-variance trade off

So can't we just choose a low-variance, low-bias modeling strategy?
Not quite!

Assume we calculate the mean squared error of some new data \mathbf{X}' given a trained model \hat{f} :

$$\text{MSE} = \mathbb{E}[(\hat{f}(\mathbf{X}') - y)^2].$$


We can decompose this further:

$$\text{MSE} = \underbrace{\mathbb{E}[(\hat{f}(\mathbf{X}') - \mathbb{E}[\hat{y}])^2]}_{\text{Variance}} + \underbrace{(\mathbb{E}[\hat{y}] - y)^2}_{\text{Bias}^2}$$

So holding the MSE fixed, if we reduce the variance we must increase the bias

- I.e. there is a **bias-variance trade-off**


Visualising the trade-off



A bit of bias can be useful


High variance

Low bias



Moderate variance

Moderate bias



Bias in ML

ML methods are typically powerful because they allow a tradeoff between variance and bias:

- ▶ We do this by “regularizing” our estimator
- ▶ Good for prediction
- ▶ Bad for inference (in simple applications)

A nice introduction to bias, regularisation and ML is provided in:
Kleinberg et al (2015). Prediction Policy Problems, AER.

Treatment effect estimation and neural networks

Effect heterogeneity

Suppose we have 8 observations of an outcome, treatment assignment and two covariates:

y	d	Gender	Education
12	1	Female	High
13	1	Female	Low
5	0	Female	High
6	0	Female	Low
7	1	Male	High
8	1	Male	Low
7	0	Male	High
6	0	Male	Low

Table 1: Observed

y	d	Gender	Education
?	0	Female	High
?	0	Female	Low
?	1	Female	High
?	1	Female	Low
?	0	Male	High
?	0	Male	Low
?	1	Male	High
?	1	Male	Low

Table 2: Unobserved counterfactual

$$ATE_{\text{Observed}} = 10 - 6 = 4$$

The ATE may mask considerable heterogeneity

Conditional Average Treatment Effect

We can break down our ATE into conditional estimates:

$$\text{CATE} = \mathbb{E}[Y|d = 1, X = a] - \mathbb{E}[Y|d = 0, X = a]$$

In particular, we can think about estimating the individual level effect, i.e.

$$\text{ITE} = [Y_i|d = 1] - [Y_i|d = 0]$$

- ▶ This is not an average
- ▶ Conventionally impossible to estimate given **fundamental problem of causal inference**

Prediction problem:

- ▶ For those treated (control) units, what would they have done under control (treatment)?

Enter the neural network

What we need is a flexible way of modelling the potential relationship between \mathbf{y} , \mathbf{d} and \mathbf{X} :

- ▶ Neural networks offer one such approach
- ▶ Very popular in industry
- ▶ Recently a lot more accessible in R
- ▶ Other methods available:
 - ▶ Random forests
 - ▶ BART
 - ▶ Best subset modelling

Simple perceptron model

Suppose we have a single vector of input data \mathbf{x} , and we want to predict the output \mathbf{y}

A simple perceptron model looks like the following:




Figure 2: Single node, single layer perceptron model

w is the weight term and b is the bias term – in this simple case, both are scalar.

Activation functions ϕ

The activation function is simply a function applied to the result of $w\mathbf{x} + b$, that controls the range of the output vector

ϕ may simply be the **identity function**:

- ▶ I.e. $\phi(\mathbf{x}) = \mathbf{x}$

Sigmoid function:


- ▶ $\phi(\mathbf{x}) = \frac{1}{1+e^{-x}}$

Rectified Linear Unit (ReLU):

- ▶ $\phi(\mathbf{x}) = \max(0, x)$

These functions (and others) are particularly useful because they have known derivatives – which we'll return to later!

Gaining a prediction from our simple model



Suppose:

- ▶ ϕ is the ReLU function
- ▶ $\mathbf{w} = \mathbf{2}, b = 1$

And we observe the following input vector \mathbf{x} :

$$\begin{bmatrix} 5 \\ 1 \\ -1 \end{bmatrix}$$

What is $\hat{\mathbf{y}}$?

Multiple inputs

The first model is very basic, so we can adapt it to accept **multiple** inputs:

- ▶ Let k index input variables, i.e. $\mathbf{X} = \{\mathbf{x}_1, \dots, \mathbf{x}_k\}$
- ▶ Let \mathbf{w} be a vector of weights, i.e. $\mathbf{w} = \{w_1, \dots, w_k\}$

Inside our activation function we replace $w\mathbf{x} + b$ with

$$w_1\mathbf{x}_1 + \dots + w_k\mathbf{x}_k + b \equiv \sum_k w_k\mathbf{x}_k + b$$




Figure 3: Single node, multiple input perceptron model

Initialisation and training

We need to set up a network structure, prior to feeding in our data.

For a single-node perceptron model with k inputs, that means instantiating the weights and biases

- ▶ A naive option sets $\mathbf{w} = \mathbf{0}$
 - ▶ This is rarely optimal – it can lead to significantly slower convergence (and can even disrupt convergence entirely)

A now standard approach is to use **Xavier initialisation** where:

$$w_k \sim \mathcal{N}(0, \frac{1}{k})$$

- ▶ where k is the number of inputs to the node
- ▶ Typically used when ϕ is tanh or sigmoidal
- ▶ Bias terms are instantiated at zero

Loss functions

The goal of the perceptron is to minimise the predictive error between \mathbf{y} and $\hat{\mathbf{y}} = \phi(\sum_k w_k \mathbf{x}_k + b)$

Depending on the type of prediction problem, we want to use a different function:

Continuous y

- ▶ ϕ will be linear or ReLU
- ▶ Minimise the mean squared error
- ▶ i.e. $\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2$

Binary y

- ▶ ϕ will be sigmoid
- ▶ Minimise using **cross-entropy** loss function
- ▶ i.e. $= -\frac{1}{N} \sum_{i=1}^n \sum_{c=1}^C y_{ic} \log(\hat{y}_{ic})$
 - ▶ where c indexes the classes within the binary/categorical variable

OLS/Logistic regression as a single-layer perceptron

We can construe OLS as a single-node perceptron model,

$$\mathbf{y} = \phi(b + w_1\mathbf{x}_1 + \dots + w_k\mathbf{x}_k),$$

when:

- ▶ ϕ is the identity function
- ▶ \mathbf{w} is the regression coefficient vector
- ▶ b is the intercept

and solved via MLE.

Similarly logistic regression is where ϕ is the sigmoid activation function.

Limitations and extensions

A single-node perceptron model is not particularly exciting:

- ▶ With identity/sigmoid activation functions we get conventional estimators
- ▶ The model is linear in inputs

To complicate our models we need to think about creating a **network** of nodes

- ▶ Increase the number of computational units
- ▶ Determine the flow of information along the network

Deep learning

Complicating the network




Figure 4: Multi-layer (but not deep) network

Deep neural network




Figure 5: Multi-layer **deep** network

Multi-layer network notation

The computation of outputs through layer h of a neural network is:

$$\mathbf{y}^{(h)} = \sigma(\mathbf{W}^{(h)}\mathbf{y}^{(h-1)} + \mathbf{b}^{(h)}),$$

where:

- ▶ $\mathbf{y}^{(h)}$ is a vector of outputs from layer h
- ▶ $\mathbf{W}^{(h)}$ is a matrix of weights for layer h
- ▶ \mathbf{b} is a vector of biases for layer h
- ▶ σ is an activation function

This model can be generalized to an arbitrary number of hidden layers H :

$$\mathbf{y} = \Phi(\mathbf{W}^{(H)}[...\sigma(\mathbf{W}^{(2)}[\sigma(\mathbf{W}^{(1)}\mathbf{x} + \mathbf{b}^{(1)})] + \mathbf{b}^{(2)})] + \mathbf{b}^{(H)}),$$

where \mathbf{x} is a vector of inputs and Φ is a final-layer activation function.

Fully-connected networks

In a fully connected network:

- Every output from layer h is an input to every node in layer $h + 1$




Figure 6: Fully-connected neural network

Feed-forward training

We initialise a multi-layer model like a single-layer model:

- ▶ Set weight terms for each node within each layer via (Xavier) initialisation

During training, an **epoch** consists of:

1. Feeding every observation through the model
 - ▶ When there are no cycles in the network, this is called “feed-forward”
2. Calculate the loss associated with the prediction
3. Adjust weights and biases based on the **gradient** of the loss
 - ▶ This is complicated with multiple layers
 - ▶ Adjusting the weights and bias affects the output of a node
 - ▶ ... and the input of the nodes (plural!) that it feeds into!
 - ▶ This process is called **backpropagation**

Estimating treatment effect heterogeneity

1. Train a neural network model on experimental data
2. Use trained model to predict counterfactual outcomes
 - ▶ Invert treatment assignment
 - ▶ Keep all covariates the same
3. Estimate ITE

$$\begin{pmatrix} \tilde{y}_{i,d=1} \\ 14 \\ 12 \\ 12 \\ 13 \\ 7 \\ 7 \\ 6 \\ 7 \end{pmatrix} - \begin{pmatrix} \tilde{y}_{i,d=0} \\ 7 \\ 7 \\ 4 \\ 6 \\ 8 \\ 6 \\ 8 \\ 6 \end{pmatrix} = \begin{pmatrix} \text{ITE} & \text{Gender} & \text{Education} \\ 7 & \text{Female} & \text{High} \\ 5 & \text{Female} & \text{Low} \\ 8 & \text{Female} & \text{High} \\ 7 & \text{Female} & \text{Low} \\ -1 & \text{Male} & \text{High} \\ 1 & \text{Male} & \text{Low} \\ -2 & \text{Male} & \text{High} \\ 1 & \text{Male} & \text{Low} \end{pmatrix}$$

4. Examine how ITE varies across covariates